Apparel Image Classification

Springboard Data Science Career Track - Capstone Two

Rose Zdybel

2020 July

Image Classification has a wide range of applications

- Medical applications
- Facial recognition
- Self driving cars image recognition
- Online retailer automatic image classification
- Product locater based on user pictures

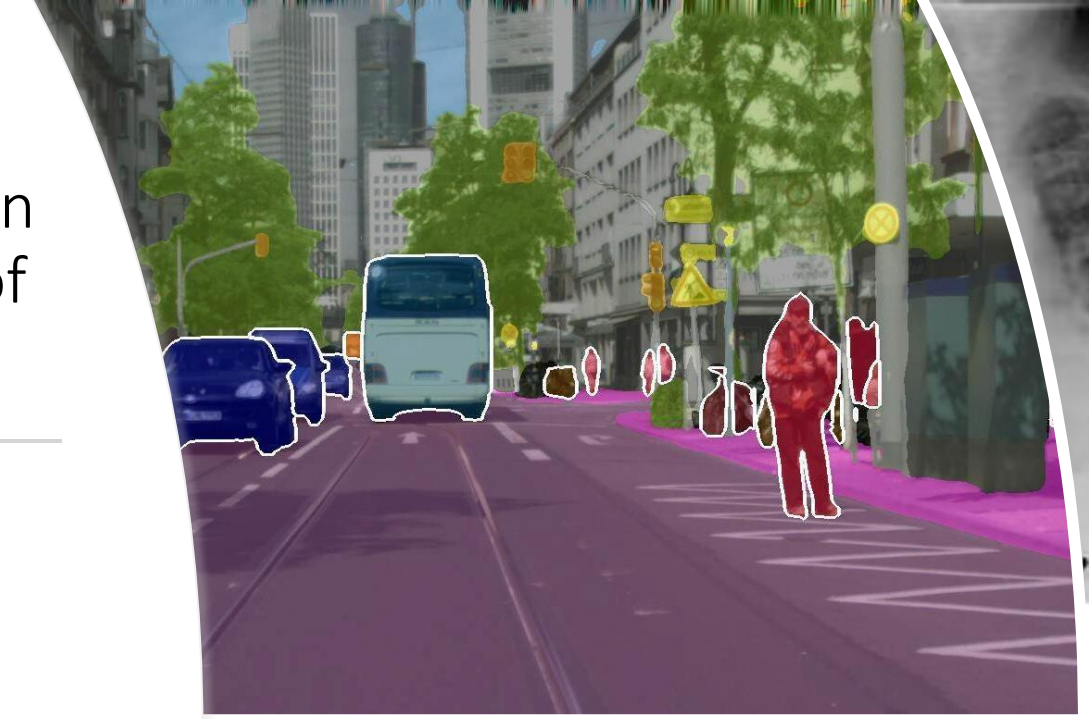

(c)

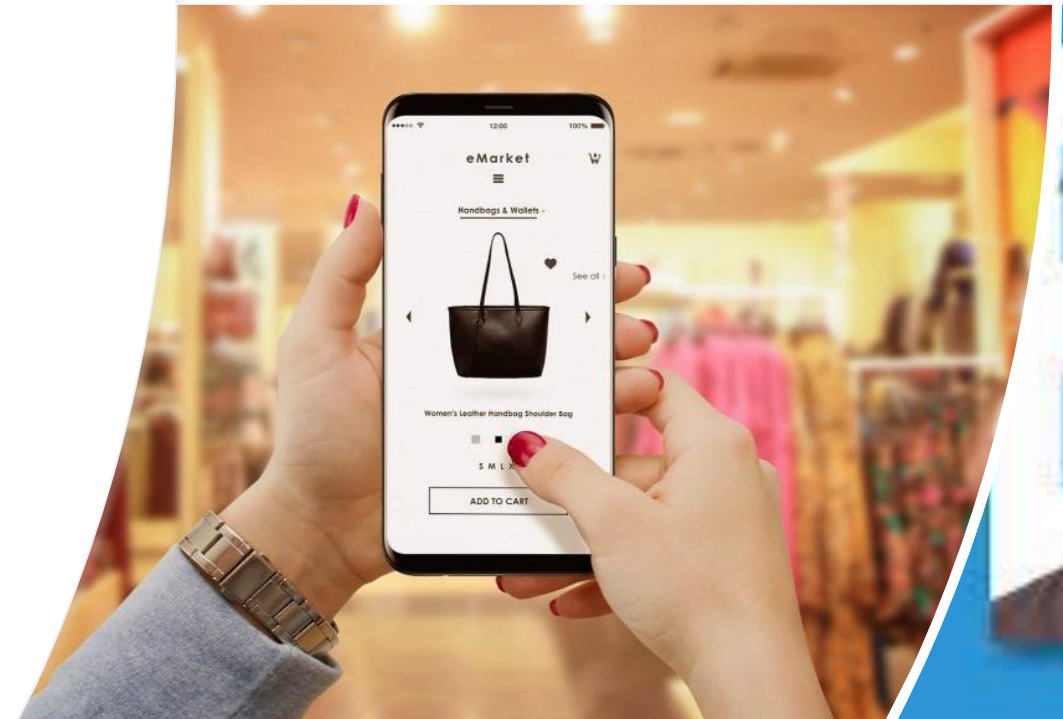

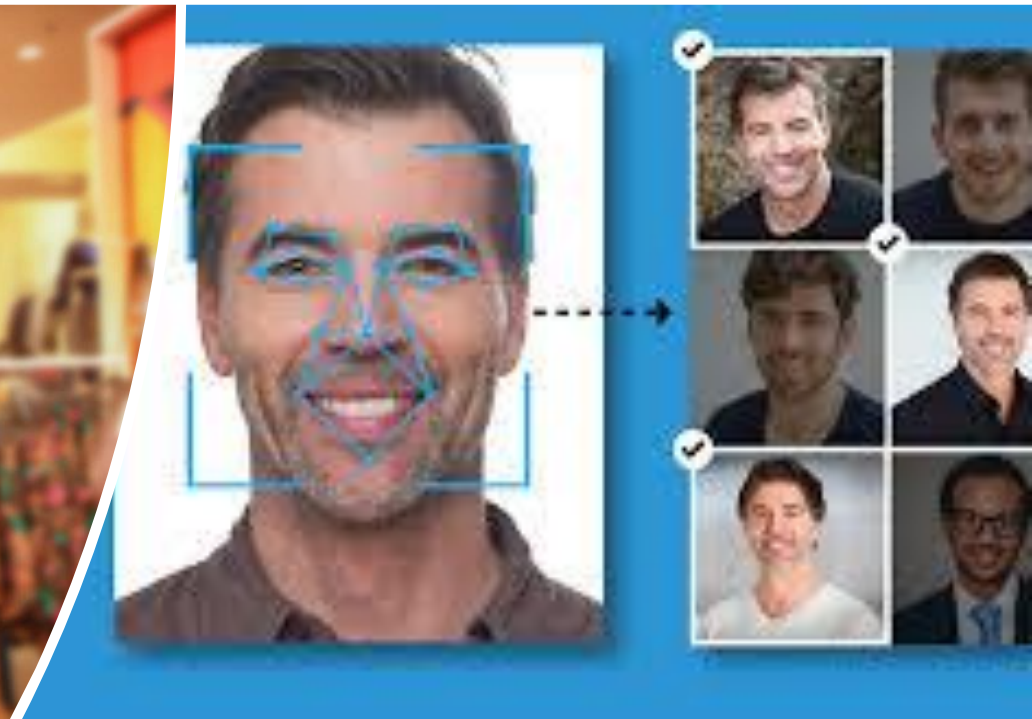

Project Objective: Image Classification Fundamentals

- Understand fundamental mechanisms used in image classification
  - To get experience in the challenges from training a model from scratch
- Maximize predictive capabilities

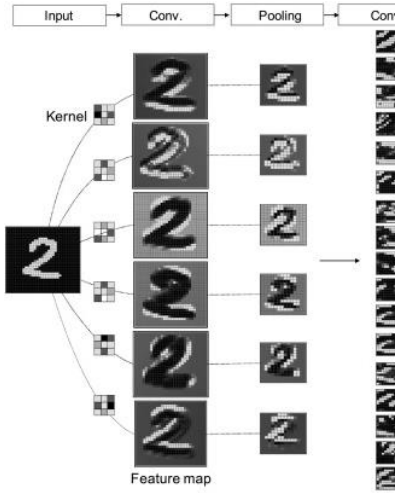

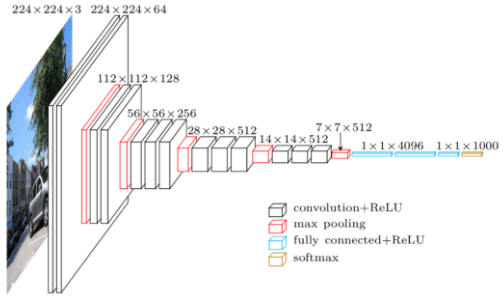

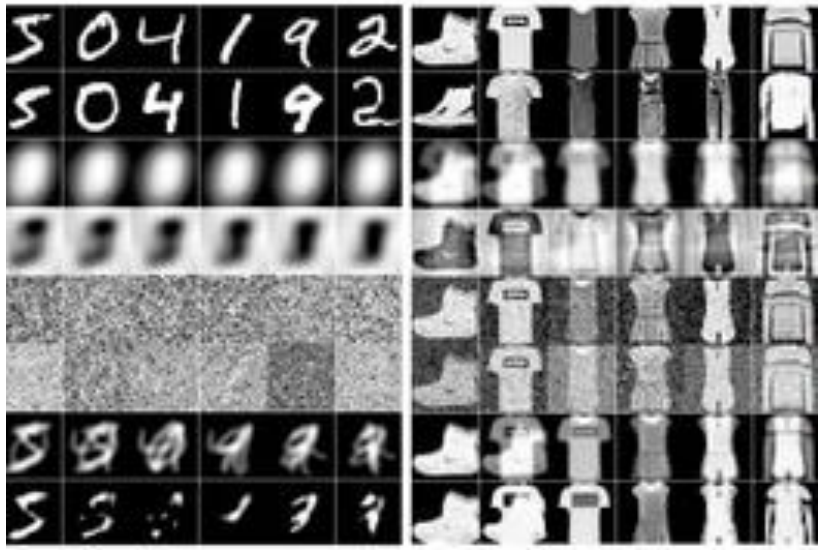

EDA

Exploratory Data Analysis

Image size

Categories

Data: Kaggle Apparel Image Data Set

- Subsetted due to resource constraints:
  - Compute:
    - Memory
    - GPU unavailability (initially)
  - Training time

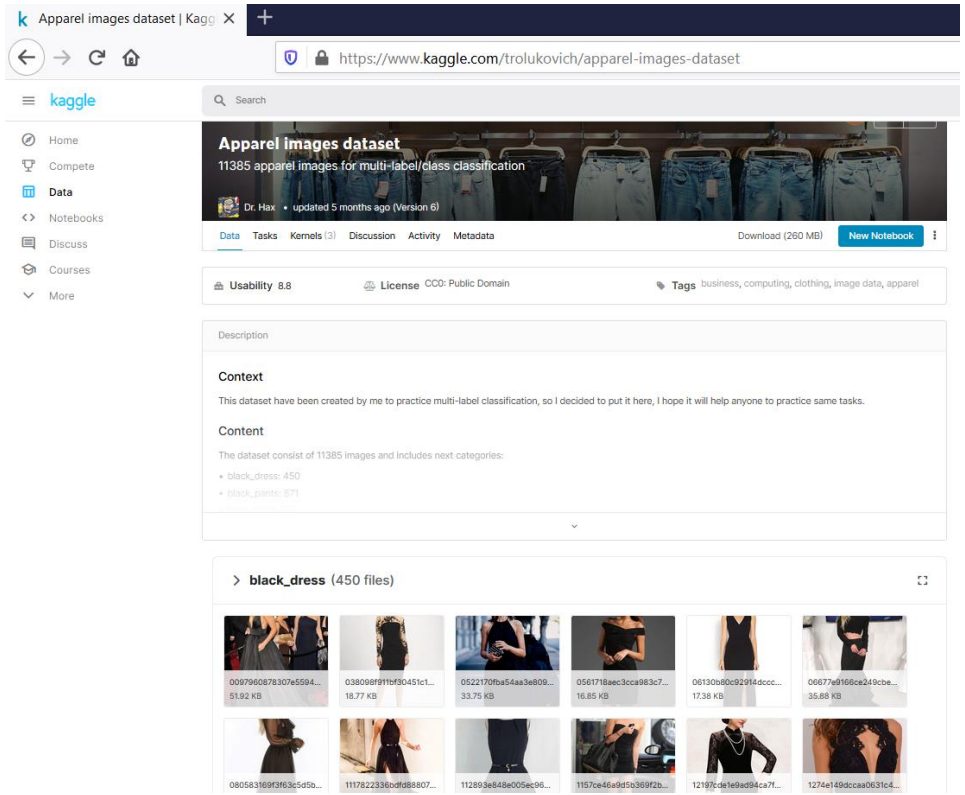

We used categories with the largest number of samples

871 black_pants
818 white_dress
800 red_dress
798 blue_pants
766 black_shoes
741 blue_shirt
715 black_shirt
610 red_shoes
600 white_shoes
523 blue_shoes
502 blue_dress
464 brown_shoes
455 green_shoes
450 black_dress
328 black_shorts
311 brown_pants
308 red_pants
299 blue_shorts
274 white_pants
230 green_shirt
227 green_pants
135 green_shorts
120 white_shorts
40 brown_shorts

Black Shirt

White Dress

Red Dress

Sample data

Approach

"Experiments"

Training	*Test*	*Validation*
0.7	0.15	0.15

Data Splits

Training

70%

Train Model Parameters

Validation

15%

Used for Early Stopping

- Loss does not decrease within specified number of Epochs

Test

15%

Measure Model Performance

Initial exploratory results influenced approach

- Issues from initial model exploration on laptop:
  - Memory deficiency
  - Training time
    - Complex models
    - Large number of pixels
  - Lack of convergence

Change	Training time	Model Stability	CPU resources	Memory resources	Accuracy
Colab Env	+		+	+	
Data set size reduction	+				-
Image size reduction options	+			+	+/-
Optimizer/LR options		+			+/-

Approach: Investigate most influential parameters

Model Architecture
- Block* configuration
- # blocks
- # nodes per layer

Optimizers & Learning Rate
- Use LR-Finder
- Verified results on selected parameter combinations

Image Size
- Source size varies
- Use single size
- Larger size increases training time

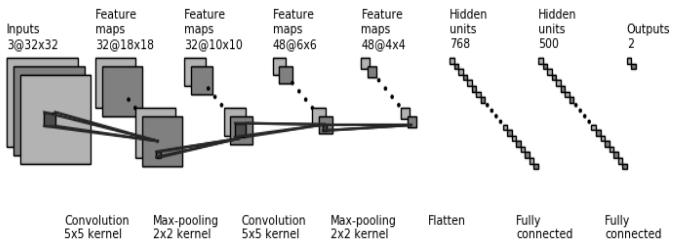

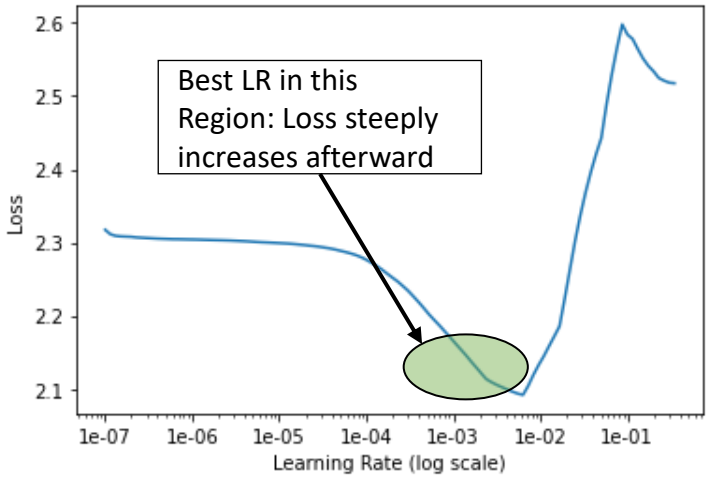

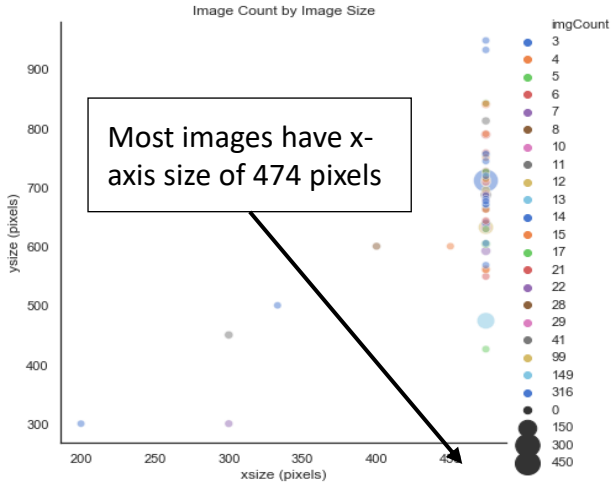

***Block:** one or more convolution layers plus a pool layer*

***Example**: Model 3; AMSGRAD optimizer:*

Parameter Grid

- Machine Learning grid search-like mechanism
  - Ran experiments with a subset of parameters based on initial exploratory results
  - Shown later in results section

OPTION	VALUES
Optimizer	SGD , AMSGRAD(AMSGrad)
Learning Rate	.01, 0.001, .0001
Models	M1, M3, M6, M7, M8
Image Size (“pixels”)	150x150, 300x300

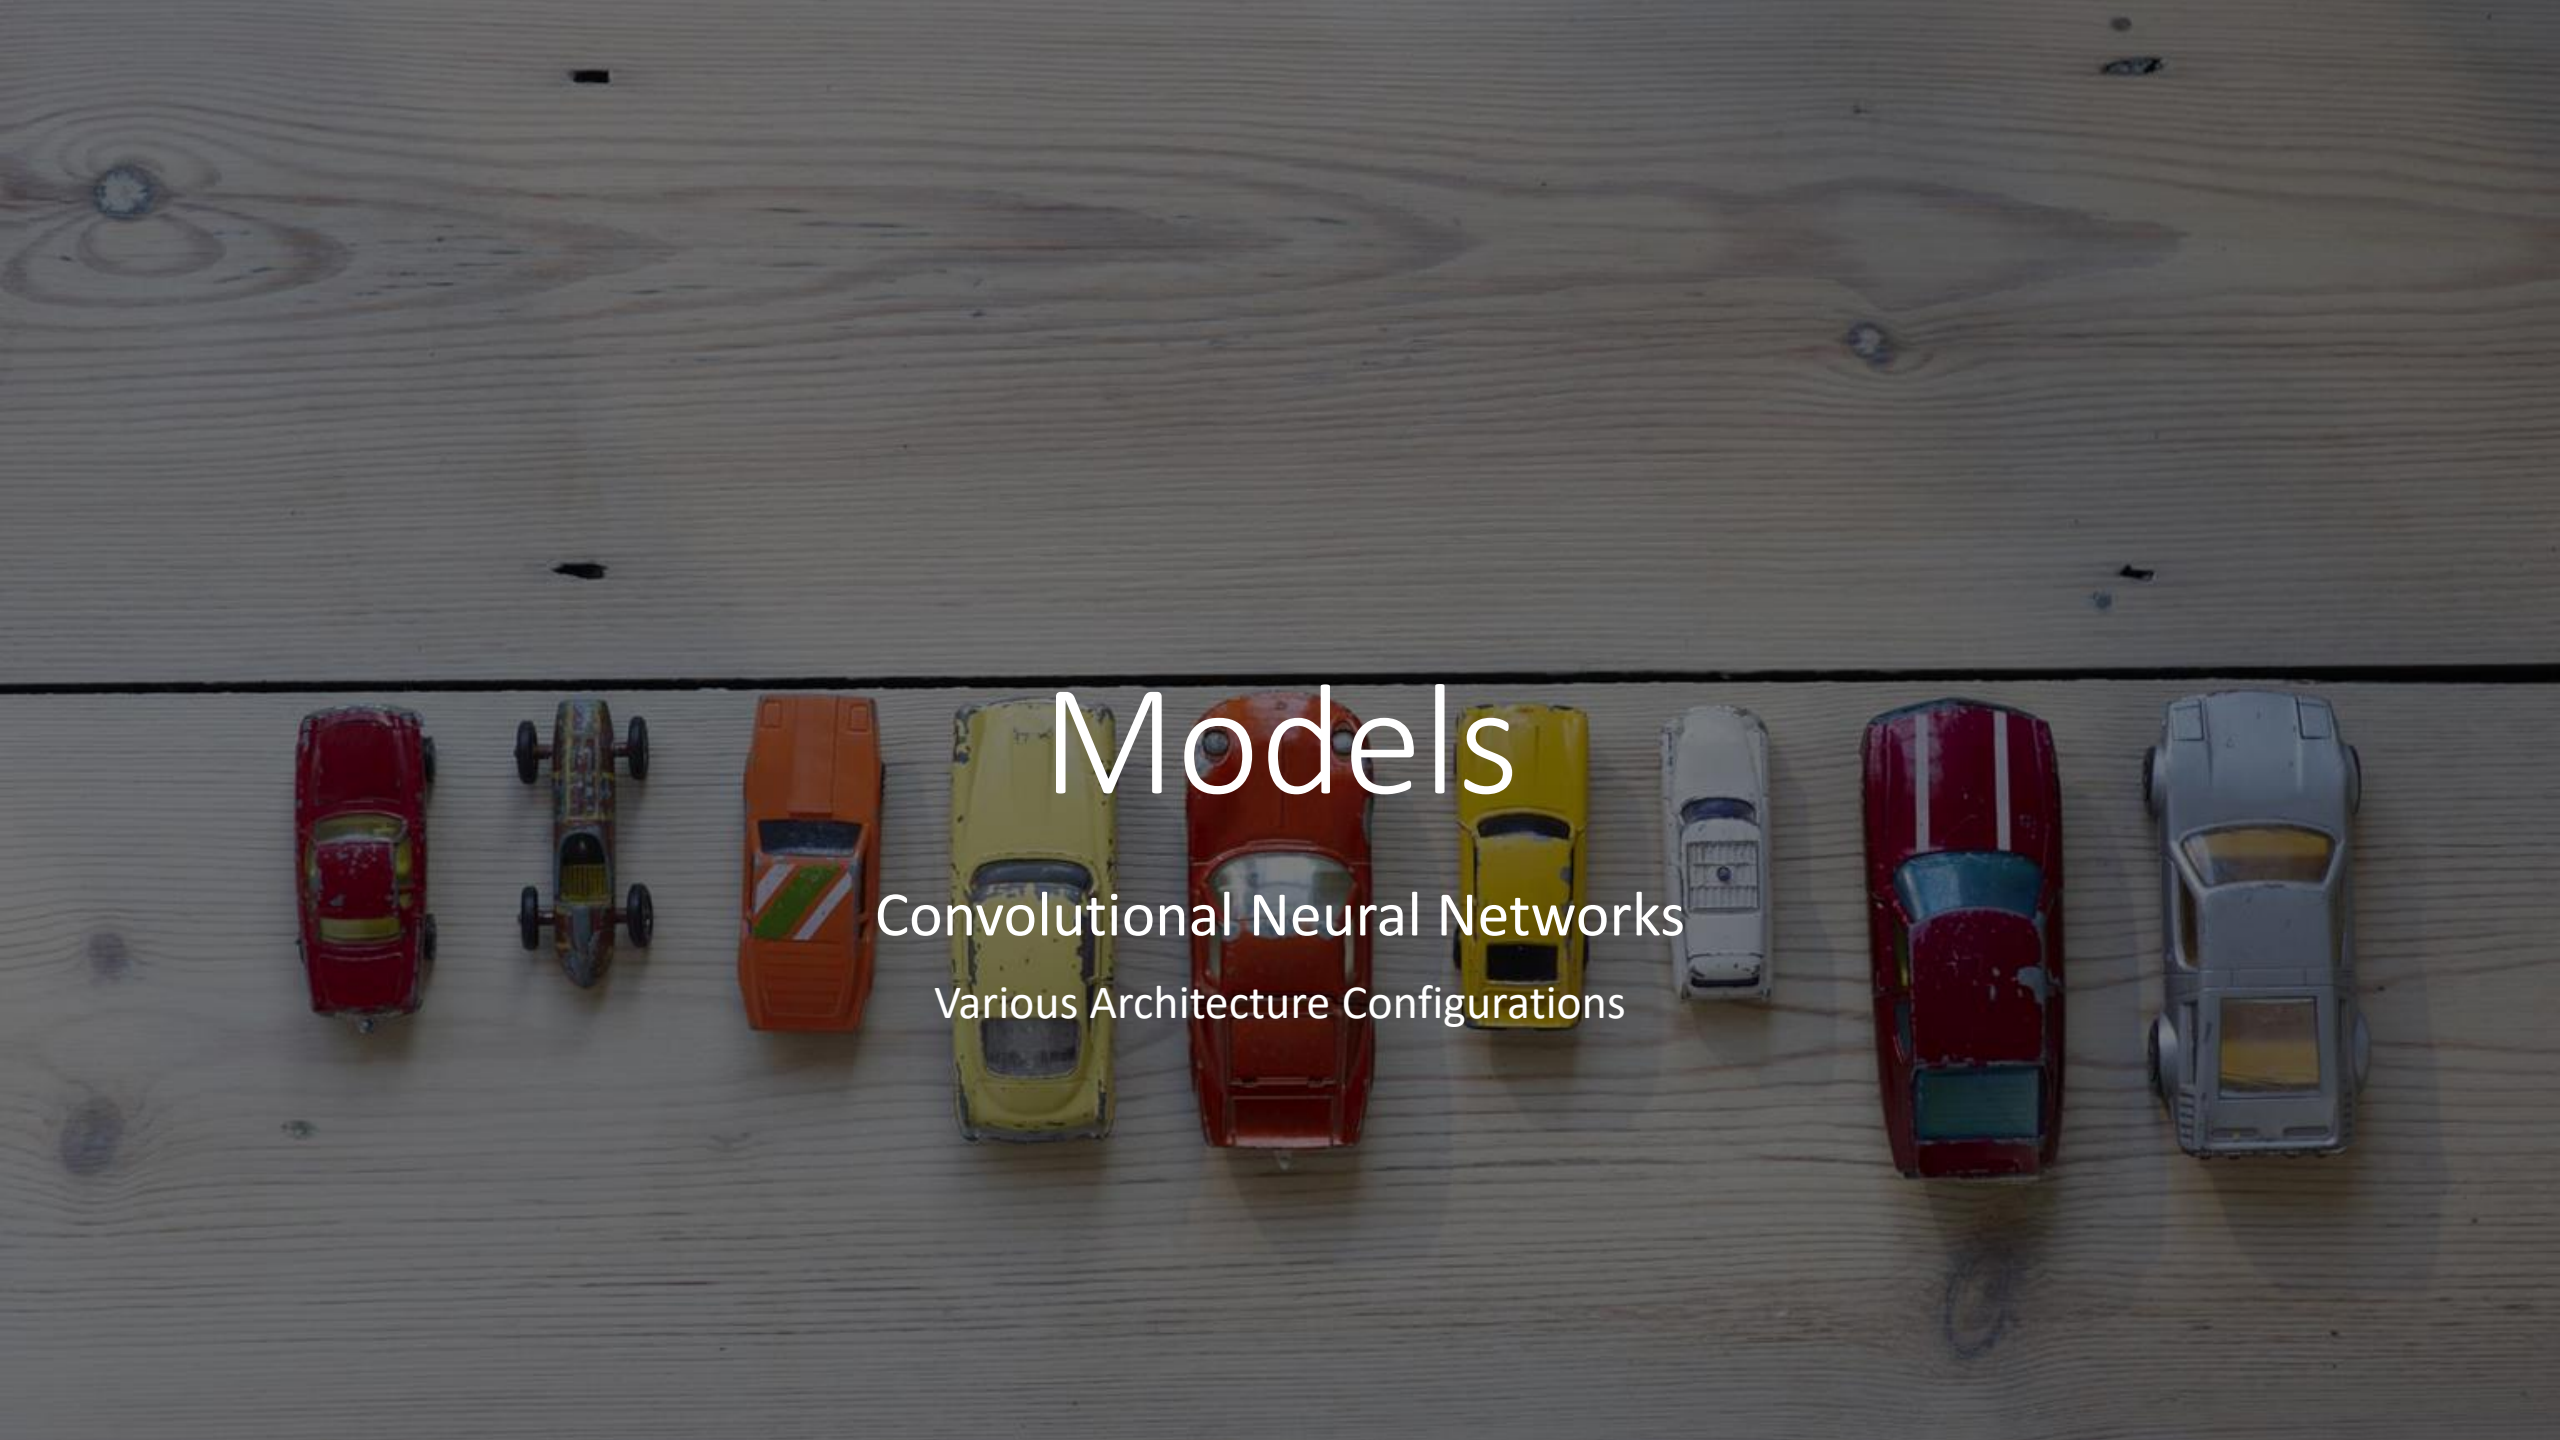

Models

Convolutional Neural Networks

Various Architecture Configurations

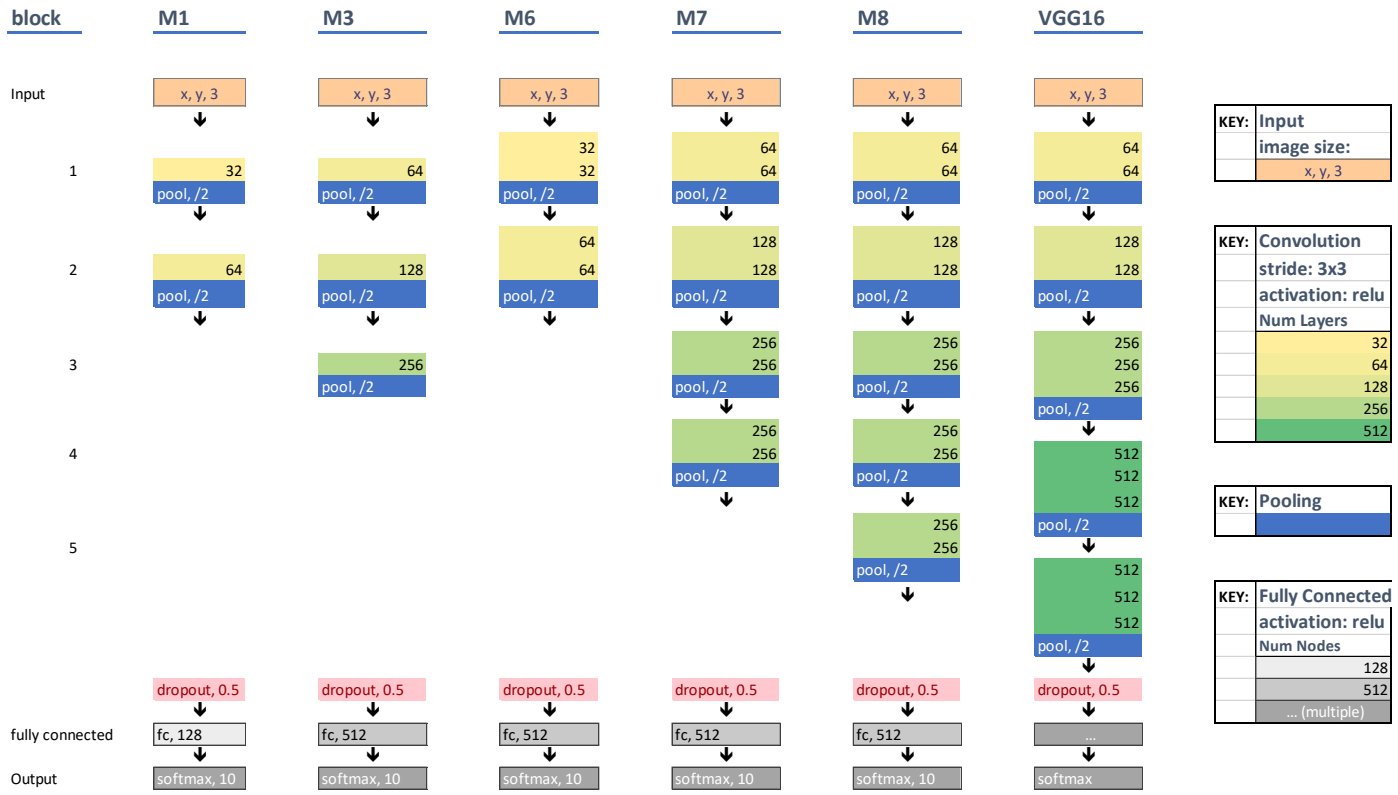

Model Architectures

- VGG16 shown for reference
  - Did not train or use pretrained model
  - For future consideration

Results

After Initial exploration, ran most promising parameter combinations

optimizer	lr	pixels	M1	M3	M6	M7	M8
Sgd	0.01	150	0.8879	0.9228	0.8796	0.9449	0.9421
Sgd	0.01	300	None	None	None	None	0.9357
Sgd	0.001	150	None	None	None	None	None
Sgd	0.001	300	None	None	None	None	None
Sgd	0.0001	150	None	None	None	None	None
Sgd	0.0001	300	None	None	None	None	None
amsgrad	0.01	150	None	None	None	None	None
amsgrad	0.01	300	None	None	None	None	None
amsgrad	0.001	150	0.9136	0.9476	0.9274	0.9449	0.9513
amsgrad	0.001	300	0.909	0.9403	0.9081	0.943	0.9338
amsgrad	0.0001	150	None	0.9292	None	0.943	None
amsgrad	0.0001	300	None	None	None	None	None

Metric: Test (Holdout) Data Accuracy

Optimizer Comparison

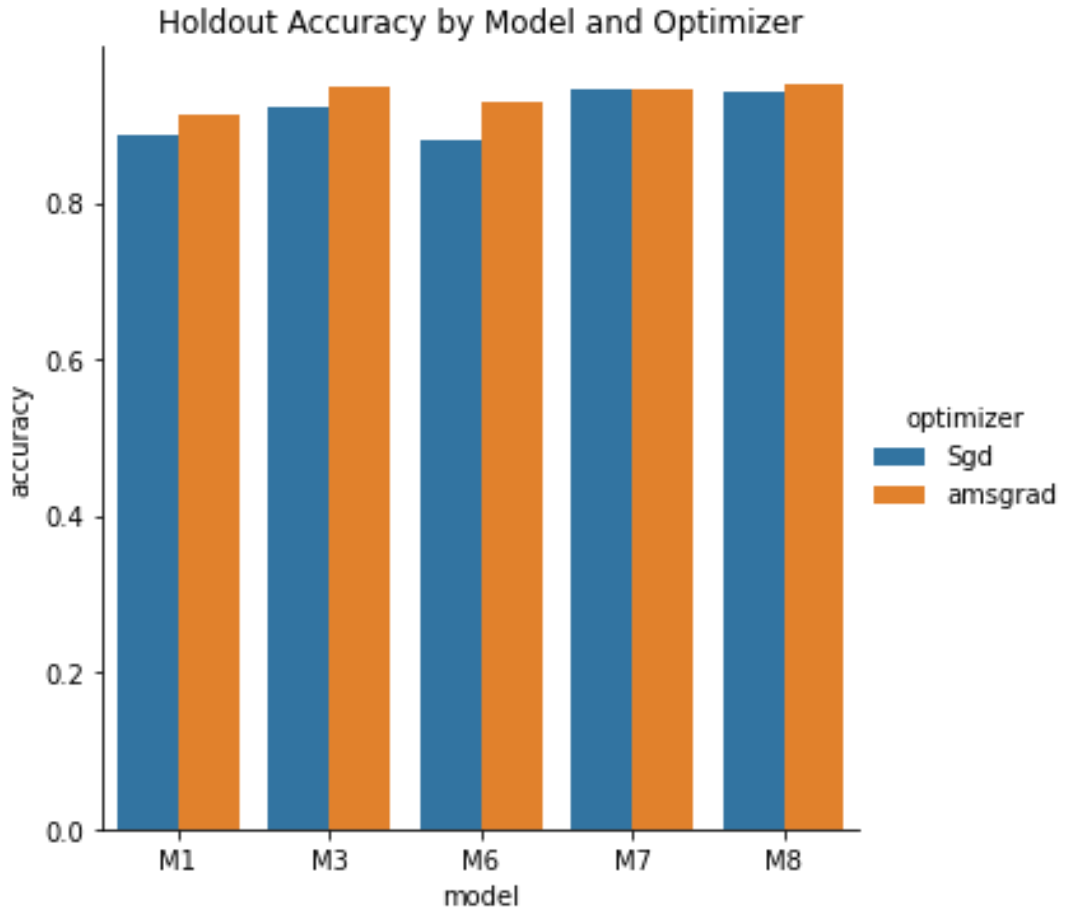

AMSGrad* did better or equivalent to SGD for all models

* **NOTE**: Adam optimizer with AMSGrad option

Image Size Comparison

- We used two sizes
  - 150 x 150
  - 300 x 300
- Smaller size produced better results
  - Should consider smaller sizes in future work

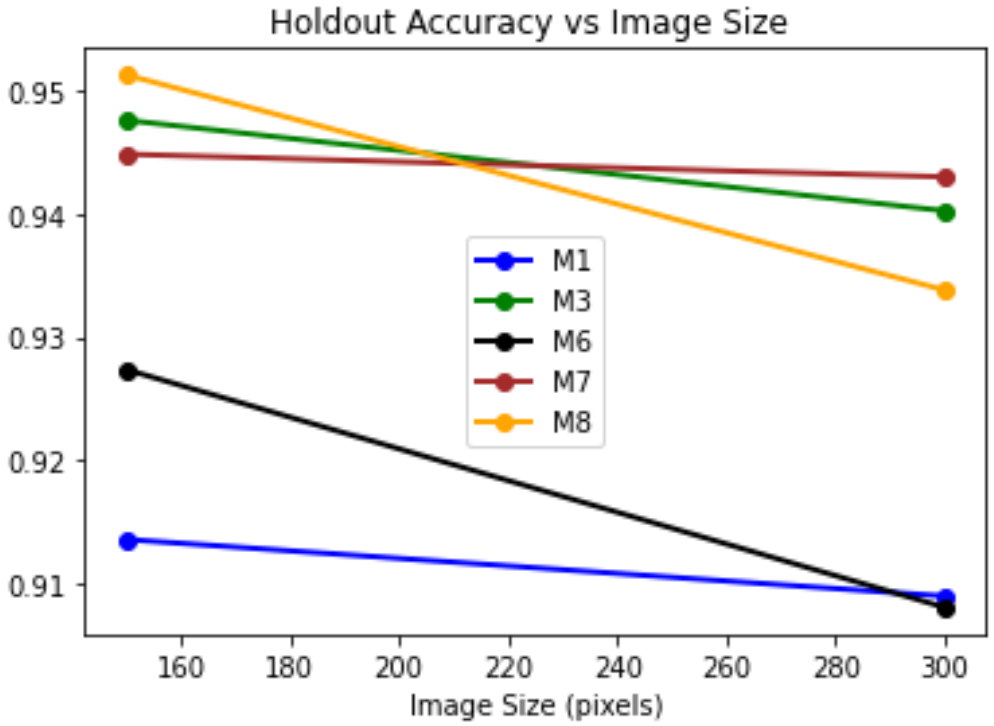

Learning Rate Comparison

- We selectively confirmed LR-finder recommendations:
  - Models
    - M3
    - M7
  - Optimizer: AMSGrad

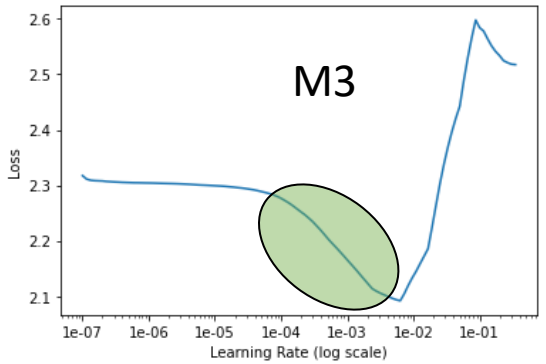

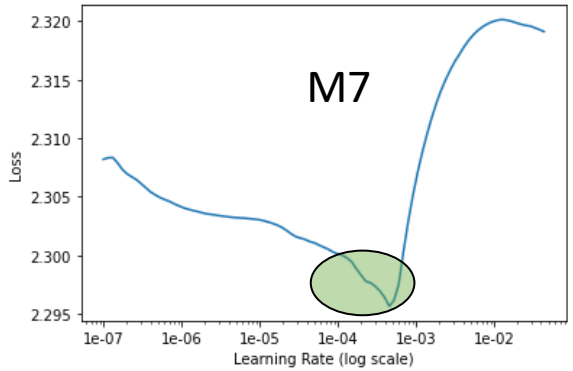

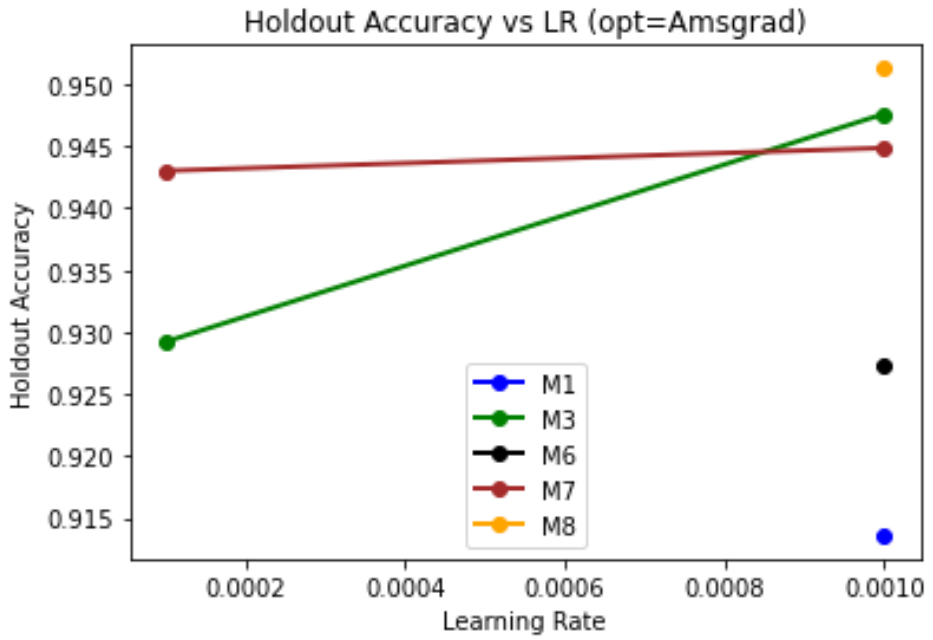

Best Model

Based on test data prediction accuracy

optimizer	lr	pixels	M1	M3	M6	M7	M8
Sgd	0.01	150	0.8879	0.9228	0.8796	0.9449	0.9421
Sgd	0.01	300	None	None	None	None	0.9357
Sgd	0.001	150	None	None	None	None	None
Sgd	0.001	300	None	None	None	None	None
Sgd	0.0001	150	None	None	None	None	None
Sgd	0.0001	300	None	None	None	None	None
amsgrad	0.01	150	None	None	None	None	None
amsgrad	0.01	300	None	None	None	None	None
amsgrad	0.001	150	0.9136	0.9476	0.9274	0.9449	0.9513
amsgrad	0.001	300	0.909	0.9403	0.9081	0.943	0.9338
amsgrad	0.0001	150	None	0.9292	None	0.943	None
amsgrad	0.0001	300	None	None	None	None	None

Test Data Accuracy	0.951
Model	M8
Image Size	150 x 150
Optimizer	AMSGrad
Learning Rate	.001

Confusion Matrix

- Major Model issues
  - White shoes classified as white dress (6)
  - Black shoes classified as black pants (5)

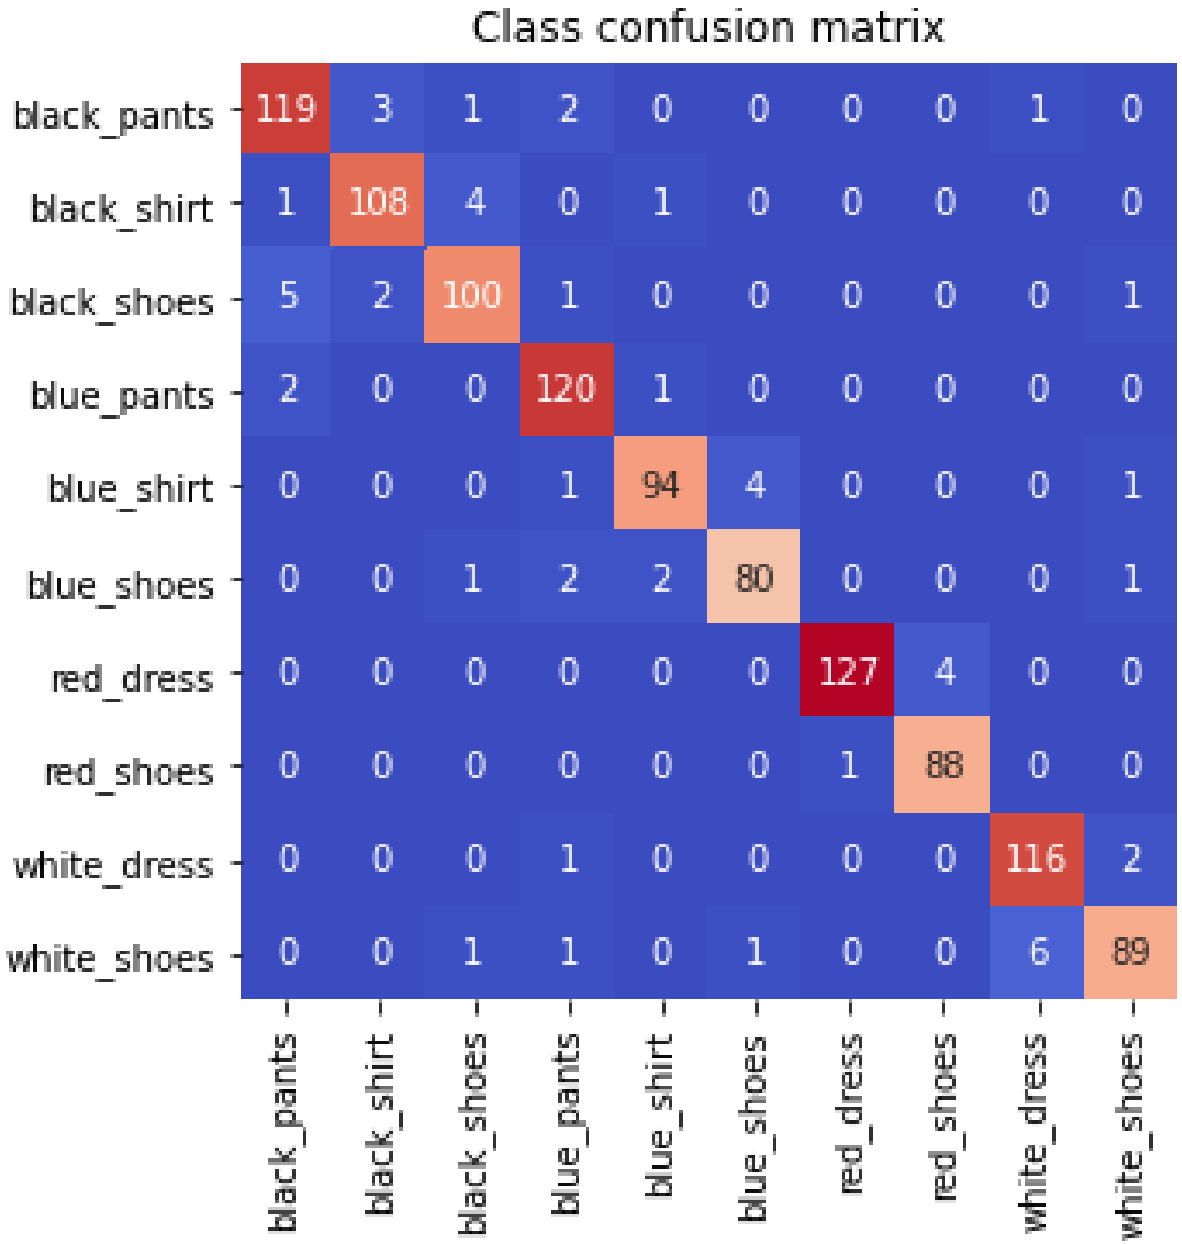

Mis-classified image examples

- Actual | Predicted
  - White shoes | white dress
  - Black shoes | black pants

Conclusion

- Explored developing a Convolution Neural Network from scratch
- Explored the effect of influential parameters
- Used Colab due to local compute resource limitations
- Developed a model with 95% predictive accuracy
-
- Future work considerations
  - Investigate multi-class model
    - Potential for better predictive capability
  - Investigate more parameter values: image sizes, optimizers
  - Investigate models with more blocks/layers
  - Investigate use of one or more pre-trained models